



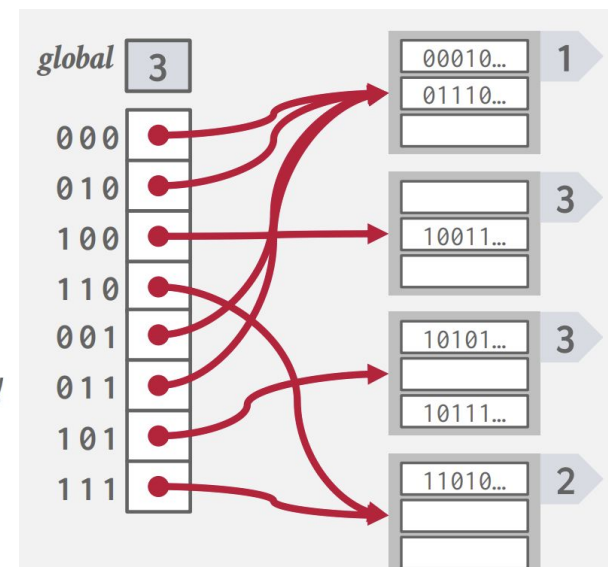


Carnegie Mellon University

Concurrent Extendible Hash Table

Shivang Dalal & Ashwin Rao
MSCS 2nd year

What is Extendible Hashing

- Ubiquitous data structure in Databases. Critical for indexes and joins.
- 2-level dynamic data structure that can perform bucket level re-hashing on overflow.
- Space efficient for sparse tables and optimized for being disk backed.
- Example of inserting a key below.



Coarse and Fine Grained locking

Coarse grained locking:

- Single directory-level read-write mutex guards all operations.
 - Correct since all hash table operations go through the directory.
- Shared lock for read operations (multiple readers simultaneously).
- Exclusive lock for write operations (ex: splitting, merging buckets).

Fine grained locking:

- Directory-level read-write mutex guards changes to directory.
- Bucket-level mutex guards data operations on a bucket.
- Early lock release strategy for contention reduction.
 - Acquire directory lock on insert/delete only for bucket split/merge.
- Global lock ordering for deadlock prevention.

Lock Free Hash Table

Lock free Bucket:

- Implemented using a lock-free linked list, kept sorted for efficiency.
- Nodes have atomic 'next' pointers updated using Compare and Swap.
- If CAS fails (due to a concurrent operation), the operation is retried in a loop until it succeeds.

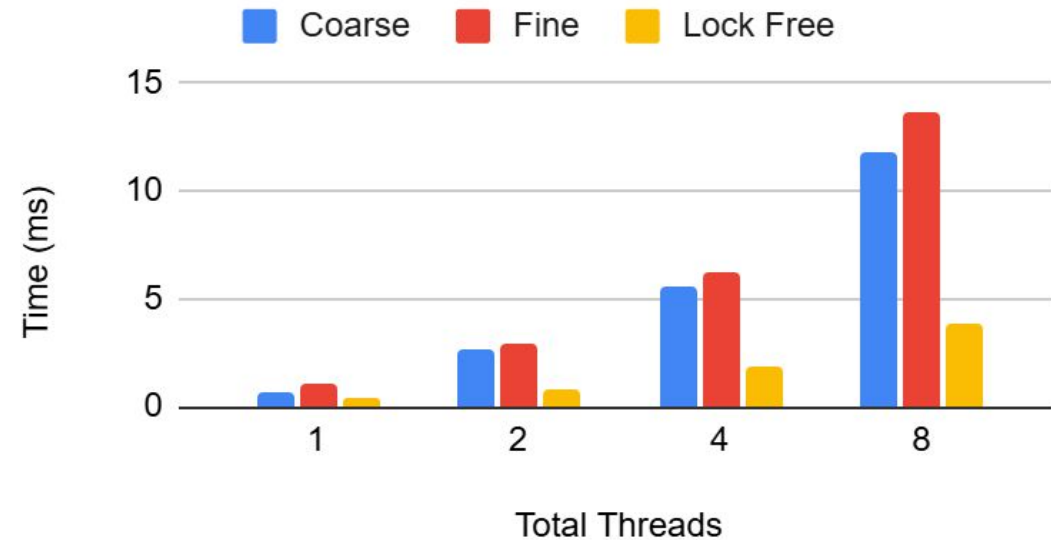
Lock free Extendible Hash Table:

- No Directory lock. Replaced with an atomic array of Bucket pointers.
- On Directory changes, the array (Bucket**) is modified using CAS.
- Memory optimization: lock-free Buckets initialized on demand.
 - CAS to ensure that each bucket is initialized by a single thread.
- No deadlocks due to no locking. Some thread always makes progress.


CRUD Benchmark

- Create, Read and Delete by every thread
- Fine grained clear favorite under high load
- Lock free better under low to medium load


CRUDTest, 100 ops per thread



CRUDTest, 1000 ops per thread



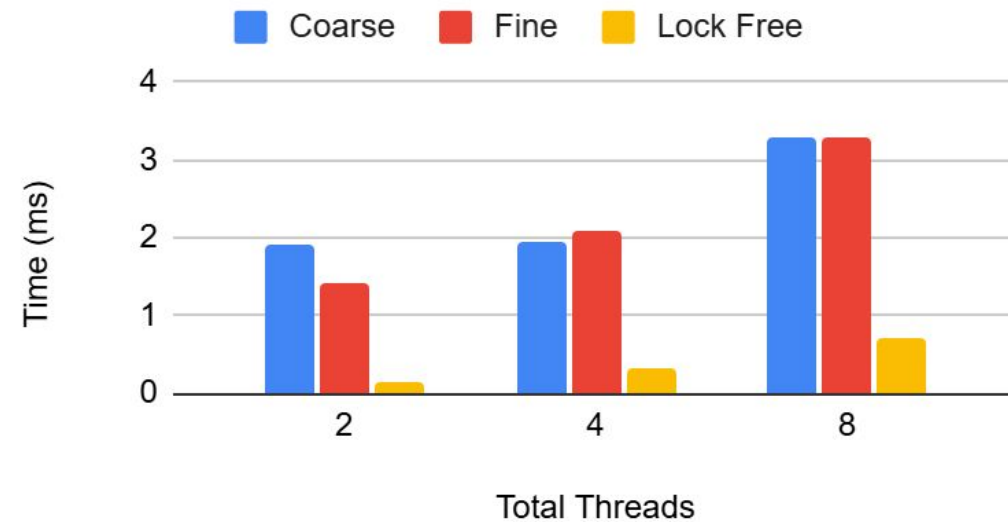
CRUDTest, 10000 ops per thread




Mixed Benchmark

- Parallel pure readers and writers. (equal ratio)
- Fine grain performs even better under high load since it can support reads more efficiently.


Mixed Test, 100 ops per thread



Mixed Test, 1000 ops per thread



Mixed Test, 10000 ops per thread



Memory and speedup

- **All three variants within 5% of each other in experiments.**
- More space efficient buckets may have seen a bigger % difference.
- Lack of critical section computation makes speedup very hard to calculate.
- Parallel always performs worse compared to serial due to overhead of sharing data and locks.
- **Artificial delays of even 10 us in worker threads showed good speedup (7-7.5x) though.**

Conclusion and recommendations

Feature	Coarse Grained	Fine Grained	Lock Free
Locking Strategy	Single directory mutex	Two-level locking	No locks. Atomic operations
Concurrency under light to medium load	Moderate	Moderate	Best
Concurrency heavy load	Bad	Best	Horrendous
Complexity	Simple	Moderate	High

- Fine-grained locking is the most suitable choice for production environments for consistent performance under varying loads.
- Lock-free implementations could be beneficial for specialized light-workload scenarios.